

# R language and data analysis: apply family

Qiang Shen

Sept.29,2016

# vectorization

- ▶ R user vs. R programmer/developer
- ▶ Vectorization is the more limited process of converting a computer program from a scalar implementation, which processes a single pair of operands at a time, to a vector implementation which processes one operation on multiple pairs of operands at once.
- ▶ Vectorization is a particular form of how parallelism is achieved.

## examples of vectorization

- ▶ vector extraction: `V[1:10]`
- ▶ vector assignment: `V[1:10] <- seq(1,10)`
- ▶ apply: `sapply(V, mean)`
- ▶ vector/matrix: `A +/- B`; `A %*% B`

# vectorization

- ▶ apply family
- ▶ plyr/reshape
- ▶ dplyr/data.table

## Looping on the Command Line

- ▶ `apply`: Apply a function over the margins of an array
- ▶ `lapply`: Loop over a list and evaluate a function on each element
- ▶ `sapply`: Same as `lapply` but try to simplify the result
- ▶ `tapply`: Apply a function over subsets of a vector

`split` is also useful, especially in conjunction with `lapply`.

# apply

- ▶ what does apply mean.
- ▶ apply to different dimension.
- ▶ compare with rowMeans etc.
- ▶ different functions.

# apply

apply is used to evaluate a function (often an anonymous one) over the **margins** of an array.

- ▶ used to apply a function to the rows or columns of a matrix or an array.
- ▶ used with general arrays, e.g. taking the average, standard deviation.
- ▶ not really faster than writing a loop, but make it simpler.

apply: margin

## Bivariate Probability Distribution

Example - Two discrete rv's  $X$  and  $Y$

| Bivariate pdf |   |     |     |     |       |
|---------------|---|-----|-----|-----|-------|
|               |   | Y   |     |     |       |
|               |   | %   | 0   | 1   | Pr(X) |
| X             | 0 | 1/8 | 0   | 1/8 |       |
|               | 1 | 2/8 | 1/8 | 3/8 |       |
|               | 2 | 1/8 | 2/8 | 3/8 |       |
|               | 3 | 0   | 1/8 | 1/8 |       |
| Pr(Y)         |   | 4/8 | 4/8 | 1   |       |

Figure 1:



# apply

```
str(apply)
```

```
## function (X, MARGIN, FUN, ...)
```

- ▶ X is an array
- ▶ MARGIN is an integer vector indicating which margins should be “retained”.
- ▶ FUN is a function to be applied
- ▶ ... is for **other arguments** to be passed to FUN

apply a function to the rows or columns of a matrix.




Figure 2: r

## apply vs. standard method

```
y<-matrix(rnorm(6),2,3)
cbind(mean(y[1,]),mean(y[2,]))
apply(y,1,mean)
```

```
x <- matrix(rnorm(30), 5, 6)
sumx<-NULL
for (i in 1:6){
  temp<-sum(x[,i])
  sumx[i]<-temp
}
sumx

apply(x,2,sum)
```

## col/row sums and means

- ▶ For sums and means of matrix dimensions, we have some shortcuts.
- ▶ `rowSums(x)`
- ▶ `rowMeans(x)`
- ▶ `colSums(x)`
- ▶ `colMeans(x)`

apply with ...

```
x <- matrix(rnorm(200), 20, 10)
apply(x, 2, quantile)
apply(x, 2, quantile, probs = c(0.25, 0.75))
```

## lapply :start from an example

- ▶ a simple function

```
##user-defined function.  
func<-function(x){  
  if (x%%2 == 0) {  
    ret<-'even'  
  }else{  
    ret<-'odd'}  
  return(ret)  
}
```

```
func(101)  
vec<-round(runif(4)*100)  
vec;func(vec)  
lapply(vec,func)
```

vectorization.

```
func<-Vectorize(func)
func(vec)
# ifelse(vec%%2, 'even', 'odd')
```

# lapply

lapply takes three arguments:

```
str(lapply)
```

```
## function (X, FUN, ...)
```



## lapply: beyond apply.

lapply always returns a list, regardless of the input.

```
lapply(iris[,1:4],mean)
```

## lapply

```
x <- list(a = 1:5, b = rnorm(10))  
x;lapply(x, mean)
```

lapply ...

```
x <- 1:4  
lapply(x, runif, min = 0, max = 10)
```

```
## [[1]]  
## [1] 7.945498  
##  
## [[2]]  
## [1] 8.216793 2.590300  
##  
## [[3]]  
## [1] 0.001757133 1.413812574 1.104137434  
##  
## [[4]]  
## [1] 1.566145 2.911776 2.376669 8.381901
```

# lapply

make use of *anonymous* functions for lapply. example: An anonymous function for extracting the 1st row of each matrix.

```
data <- list(a = matrix(1:6, 2, 3), b = matrix(1:6, 3, 2),  
data  
lapply(data, function(x) x[1,])
```

# lapply

*anonymous* functions continued.

```
lapply(iris[,1:4],function(x) sd(x,na.rm=T)/mean(x,na.rm=T))
myfunc<-function(x){
  rec<-c(mean(x,na.rm=T),sd(x,na.rm=T))
  return(rec)
}
result<-lapply(iris[,1:4],myfunc)
result
```

methods to covert list into data.frame.

```
t(as.data.frame(result))  
# t(sapply(result, '['))  
do.call('rbind',result)
```

# sapply

sapply will try to simplify the result of lapply if possible.

- ▶ If the result is a list where every element is length 1, then a vector is returned
- ▶ If the result is a list where every element is a vector of the same length ( $> 1$ ), a matrix is returned.
- ▶ If it can't figure things out, a list is returned

## tapply: split-apply-combine.

- ▶ Split up a big dataset
- ▶ Apply a function to each piece
- ▶ Combine all the pieces back together
- ▶ map-reduce in hadoop.




Figure 3:



## tapply




tapply is used to apply a function over subsets of a vector.

```
str(tapply)
```

```
## function (X, INDEX, FUN = NULL, ..., simplify = TRUE)
```

## tapply

Take group means.

|       |             | split | Apply   | combine   |
|-------|-------------|-------|---|-----------|
|       |             | x f   |   |           |
| [1,]  | -0.56047565 | 1     |  | 0.1935703 |
| [2,]  | -0.23017749 | 1     |   |           |
| [3,]  | 1.55870831  | 1     |   |           |
| [4,]  | 0.07050839  | 1     |   |           |
| [5,]  | 0.12928774  | 1     |   |           |
| [6,]  | 0.95683335  | 2     |  | 0.5526592 |
| [7,]  | 0.45333416  | 2     |   |           |
| [8,]  | 0.67757064  | 2     |   |           |
| [9,]  | 0.57263340  | 2     |   |           |
| [10,] | 0.10292468  | 2     |   |           |
| [11,] | 2.28055488  | 3     |  | 1.8551235 |
| [12,] | -0.72727063 | 3     |   |           |
| [13,] | 2.69018435  | 3     |   |           |
| [14,] | 1.50381245  | 3     |   |           |
| [15,] | 3.52833655  | 3     |   |           |

## tapply

```
set.seed(123)
x <- c(rnorm(5), runif(5), rnorm(5, 1))
f <- gl(3, 5)
data<-cbind(x,f)
tapply(x, f, mean)
```

```
##           1           2           3
## 0.1935703 0.5526592 1.8551235
```

```
tapply(x, f, mean, simplify = FALSE)
```

```
## $`1`
## [1] 0.1935703
##
## $`2`
## [1] 0.5526592
##
## $`3`
```

## tapply

Find group ranges.

```
tapply(x, f, range)
```

```
## $`1`
```

```
## [1] -0.5604756  1.5587083
```

```
##
```

```
## $`2`
```

```
## [1] 0.1029247 0.9568333
```

```
##
```

```
## $`3`
```

```
## [1] -0.7272706  3.5283366
```

## tapply example

```
tapply(iris[,1], iris$Species, mean)
```

```
##      setosa versicolor  virginica  
##      5.006      5.936      6.588
```

# split

`split` takes a vector or other objects and splits it into groups determined by a factor or list of factors.

```
str(split)
```

```
## function (x, f, drop = FALSE, ...)
```

```
# function (x, f, drop = FALSE, ...)
```

- ▶ `x` is a vector (or list) or data frame
- ▶ `f` is a factor (or coerced to one) or a list of factors
- ▶ `drop` indicates whether empty factors levels should be dropped

## split

```
x <- c(rnorm(10), runif(10), rnorm(10, 1))  
f <- gl(3, 10)  
split(x, f)
```

```
## $`1`
```

```
## [1] 0.54909674 0.23821292 -1.04889314 1.29476325 0
```

```
## [6] -0.05568601 -0.78438222 -0.73350322 -0.21586539 -0
```

```
##
```

```
## $`2`
```

```
## [1] 0.13880606 0.23303410 0.46596245 0.26597264 0.85782
```

```
## [7] 0.44220007 0.79892485 0.12189926 0.56094798
```

```
##
```

```
## $`3`
```

```
## [1] 0.1814843 1.6849361 0.6799436 -0.3115224 0.4003
```

```
## [7] 1.8867361 0.8486040 1.3297912 -2.2273228
```

## split

A common idiom is `split` followed by an `lapply`.

```
lapply(split(x, f), mean)
```

```
## $`1`  
## [1] -0.026563  
##  
## $`2`  
## [1] 0.3931406  
##  
## $`3`  
## [1] 0.5343631
```



## Splitting a Data Frame

```
s <- split(iris, iris$Species)
sapply(s, function(x) colMeans(x[, 1:4], na.rm=T))
```

```
##           setosa versicolor virginica
## Sepal.Length  5.006         5.936         6.588
## Sepal.Width   3.428         2.770         2.974
## Petal.Length  1.462         4.260         5.552
## Petal.Width   0.246         1.326         2.026
```

```
# sapply(s, function(x) sapply((x[, 1:4]), mean, na.rm=T))
```

## Splitting on More than One Level

```
x <- rnorm(10)
f1 <- gl(2, 5)
f2 <- gl(5, 2)
f1;f2
```

```
## [1] 1 1 1 1 1 2 2 2 2 2
```

```
## Levels: 1 2
```

```
## [1] 1 1 2 2 3 3 4 4 5 5
```

```
## Levels: 1 2 3 4 5
```

```
interaction(f1, f2)
```

```
## [1] 1.1 1.1 1.2 1.2 1.3 2.3 2.4 2.4 2.5 2.5
```

```
## Levels: 1.1 2.1 1.2 2.2 1.3 2.3 1.4 2.4 1.5 2.5
```

```
data<-cbind(x,f1,f2)
```

```
data
```

```
##           x         f1         f2
```

split:Empty levels can be dropped.

```
str(split(x, list(f1, f2), drop = TRUE))
```

```
## List of 6
## $ 1.1: num [1:2] -0.772 0.287
## $ 1.2: num [1:2] -1.221 0.435
## $ 1.3: num 0.8
## $ 2.3: num -0.164
## $ 2.4: num [1:2] 1.243 -0.934
## $ 2.5: num [1:2] 0.394 0.404
```

# Looping on the Command Line

- ▶ `lapply`: Loop over a list and evaluate a function on each element
- ▶ `sapply`: Same as `lapply` but try to simplify the result
- ▶ `apply`: Apply a function over the margins of an array
- ▶ `tapply`: Apply a function over subsets of a vector